

Tin Htet Kyaw

Hands-on Software Engineering Leader

Millbrae, CA, USA

📞 +1 (650) 250-0888 | ✉️ tinhkyaw@alum.mit.edu | 📱 tinhkyaw | 🌐 tinhkyaw

Hands-on software engineering leader with experience in building distributed, scalable Big Data, Machine Learning, Artificial Intelligence systems

Experience

Lightspeed Venture Partners, Menlo Park, CA

Vice President, Data

Jan. 2022 - Present

- Leading the efforts in scaling and driving the adoption of data platform at Lightspeed
 - Lightspeed's data platform tracks over 3.8 million companies and 270 million people to search for investment opportunities
 - Our automated data room diligence tool reduces the time to perform due diligence on a company from days to less than 30 minutes
 - Lightspeed data team was voted as one of the top 20 most data driven venture capital funds by peers in Data Driven VC Landscape Report 2 years in a row for 2023 and 2024

Director of Engineering/Founding Member of Lightspeed's Data Team

Jan. 2019 - December 2021

- Led the Data Engineering team at Lightspeed to create a data platform to power deal sourcing, diligence and portfolio support in venture capital investing
 - Built scalable data infrastructure to onboard datasets, generate features, train and test machine learning models
 - Created internal tools to help track and analyze relationships and signals on future investment opportunities and candidates
 - Worked cross-functionally with investors and operations teams to improve the quality of data models and automate business processes

Lyft Inc., San Francisco, CA

Machine Learning & Intelligence Staff Software Engineer, Autonomous Vehicle Division

Jun. 2018 - Jan. 2019

- Built a Kubernetes-based machine learning pipeline to easily iterate on Tensorflow and PyTorch computer vision models to build High Definition maps

Staff Software Engineer/Engineering Lead

Aug. 2017 - May 2018

- Led a team of three engineers to rebuild Lyft's live traffic profile estimation pipeline using PySpark, reducing its overall latency from hours to less than 30 minutes — the pipeline produced traffic models simulating real-time traffic conditions based on recently-live Lyft rides and served as a critical input to Lyft's ETA estimation models
- Collaborated with ETA Data Science team to reduce Lyft's ETA estimation error rate by 5%

Walmart Labs, Sunnyvale, CA

Principal Software Engineer/Engineering Lead

Aug. 2016 - Apr. 2017

- Led a team of three engineers to re-architect the Walmart Ad Exchange (WMX) platform on Apache Spark to drive smart and responsive ad-targeting based on Walmart's unique offline+online datasets

Netflix Inc., Los Gatos, CA

Senior Software Engineer

Aug. 2014 - Jul. 2016

- Built and launched the first production-grade machine learning pipeline on Apache Spark at Netflix to:
 - Enable rapid development of recommendation algorithms for movies and TV shows
 - Generate daily personalized recommendations to power email, push notifications and Billboard — the large banner above the fold on Netflix homepage
- Wrote the recommendation algorithm to cold-start new titles with little metadata and no play history
- Improved Billboard recommendation algorithm to balance impressions among competing titles while maximizing streaming hours

Drawbridge Inc. (acquired by LinkedIn), San Mateo, CA

First Engineer, Algorithm Engineering Manager

Jan. 2011 - Feb. 2014

- Created Drawbridge's first Device Graph to statistically connect anonymized ids of the same user across multiple devices
 - Achieved over 60% precision on a partner-provided test dataset
- Designed and built Drawbridge's batch data pipeline using Hadoop, Hive, Pig and Oozie to process daily billions of events in terabytes of data
- Managed a team of three engineers building machine learning models for:
 - Click-through rate (CTR) and conversion rate (CVR) prediction and optimization
 - Loss prevention — to flag, stop and alert for ad groups and campaigns where losses, differences between predicted and actual revenues, exceed alert thresholds
 - User behavior, demographic and geographic feature extraction and extrapolation

Announce Media Inc. (Answers.com), Mountain View, CA

Senior Software Engineer

Jun. 2010 - Dec. 2010

- Architected and launched Announce Media's first Hadoop application to parallelize, scale and automate keyword evaluation for the company's core business of search engine marketing

AdMob Inc. (acquired by Google), San Mateo, CA

Software Engineer

May 2008 - May 2010

- Built and maintained Unique User (UU) service, a real-time system keeping track of ad impressions and clicks per UU throughout the hyper-growth phase of the early mobile web
- Created Hadoop jobs to compute analytics of UU data in AdMob's advertising network

Yahoo! Inc., Sunnyvale, CA

Software Engineer

Sep. 2005 - Apr 2008

- Developed ad-serving, link-tracking, A/B testing, location-based personalization and backend for core modules such as Mail, RSS, Scoreboard, TV, Weather and Personal Assistant modules to My Yahoo!

Education

Massachusetts Institute of Technology, Cambridge, MA

Sep. 2000 - Sep. 2005

- Master of Engineering in Electrical Engineering and Computer Science (EECS) *Sep. 2005*
- Bachelor of Science in Management Science with concentration in Finance *Jun. 2005*
- Bachelor of Science in Electrical Engineering and Computer Science (EECS) with Minor in Economics *Jun. 2004*

Skills

Languages & Frameworks: Java, Scala, Python, Bash, SQL, Spark, Hadoop, Hive, Airflow, Docker, Kubernetes, PostgreSQL, AWS, Git
Communication: English, Burmese, Chinese Mandarin, Japanese (3 years of college-level coursework)

Patents & Awards

Method for Transmitting Information to Device on a Network Based upon Information Associated with a Different Device on the Network, Sivaramakrishnan, K.; Guan, D.; Kyaw T.; et al. 2016.

US 20160125461 A1

System and Method for Determining Related Digital Identities, Sivaramakrishnan, K.; Guan, D.; Kyaw T.; et al. 2014.

US 20140095320 A1

Fourth Place, National Matriculation Examination, Ministry of Education, Republic of the Union of Myanmar

1997