Tin Htet Kyaw
Hands-on Software Engineering Leader
Millbrae, CA, USA

📱 (650) 618-8168 | ✉ tinhkyaw@alum.mit.edu | ⌨ tinhkyaw | 🔗 tinhkyaw

Hands-on engineering leader with wealth of experience building distributed, scalable Big Data, Machine Learning systems

Experience

Lightspeed Venture Partners, *Menlo Park, CA*

Director of Engineering/Founding Member of Lightspeed's Data Platform Team   *Jan. 2019 - Present*

- Leading a team of two engineers and five contractors to create automated, data-driven solutions for deal sourcing, diligence and portfolio support in venture capital investing
    - Building scalable data infrastructure to onboard datasets, generate features, train and test machine learning models
    - Creating internal tools to help track and analyze relationships and signals on future investment opportunities and candidates
    - Working cross-functionally with investors to improve the ability of our data models to generate useful business insights

Lyft Inc., *San Francisco, CA*

Machine Learning & Intelligence Staff Software Engineer, Autonomous Vehicle Division   *Jun. 2018 - Jan. 2019*

- Built a Kubernetes-based machine learning pipeline to easily iterate on Tensorflow and PyTorch computer vision models to build High Definition maps

Staff Software Engineer/Engineering Lead   *Aug. 2017 - May 2018*

- Led a team of three engineers to rebuild Lyft's live traffic profile estimation pipeline using PySpark, reducing its overall latency from hours to less than 30 minutes — the pipeline produced traffic models simulating real-time traffic conditions based on recently-live Lyft rides and served as a critical input to Lyft's ETA estimation models
- Collaborated with ETA Data Science team to reduce Lyft's ETA estimation error rate by 5%

Walmart Labs, *Sunnyvale, CA*

Principal Software Engineer/Engineering Lead   *Aug. 2016 - Apr. 2017*

- Led a team of three engineers to re-architect the Walmart Ad Exchange (WMX) platform on Apache Spark to drive smart and responsive ad-targeting based on Walmart's unique offline+online datasets

Netflix Inc., *Los Gatos, CA*

Senior Software Engineer   *Aug. 2014 - Jul. 2016*

- Built and launched the first production-grade machine learning pipeline on Apache Spark at Netflix to:
    - Enable rapid development of recommendation algorithms for movies and TV shows
    - Generate daily personalized recommendations to power email, push notifications and Billboard — the large banner above the fold on Netflix homepage
- Wrote the recommendation algorithm to cold-start new titles with little metadata and no play history
- Improved Billboard recommendation algorithm to balance impressions among competing titles while maximizing streaming hours

Drawbridge Inc. *(acquired by LinkedIn)*, *San Mateo, CA*

First Engineer, Algorithm Engineering Manager   *Jan. 2011 - Feb. 2014*

- Created Drawbridge's first Device Graph to statistically connect anonymized ids of the same user across multiple devices
    - Achieved over 60% precision on a partner-provided test dataset
- Designed and built Drawbridge's batch data pipeline using Hadoop, Hive, Pig and Oozie to process daily billions of events in terabytes of data
- Managed a team of three engineers building machine learning models for:
    - Click-through rate (CTR) and conversion rate (CVR) prediction and optimization
    - Loss prevention — to flag, stop and alert for ad groups and campaigns where losses, differences between predicted and actual revenues, exceed alert thresholds
    - User behavior, demographic and geographic feature extraction and extrapolation

Announce Media Inc. *(Answers.com)*, *Mountain View, CA*

Senior Software Engineer   *Jun. 2010 - Dec. 2010*

- Architected and launched Announce Media's first Hadoop application to parallelize, scale and automate keyword evaluation for the company's core business of search engine marketing

AdMob Inc. *(acquired by Google)*, *San Mateo, CA*

Software Engineer *May 2008 - May 2010*
- Built and maintained Unique User (UU) service, a real-time system keeping track of ad impressions and clicks per UU throughout the hyper-growth phase of the early mobile web
- Created Hadoop jobs to compute analytics of UU data in AdMob's advertising network

Yahoo! Inc., *Sunnyvale, CA*

Software Engineer *Sep. 2005 - Apr 2008*
- Developed ad-serving, link-tracking, A/B testing, location-based personalization and backend for core modules such as Mail, RSS, Scoreboard, TV, Weather and Personal Assistant modules to My Yahoo!

Education

Massachusetts Institute of Technology, *Cambridge, MA* *Sep. 2000 - Sep. 2005*
- Master of Engineering in Electrical Engineering and Computer Science (EECS) *Sep. 2005*
- Bachelor of Science in Management Science with concentration in Finance *Jun. 2005*
- Bachelor of Science in Electrical Engineering and Computer Science (EECS) with Minor in Economics *Jun. 2004*

Skills

Languages & Frameworks: Java, Scala, Python, Bash, SQL, Spark, Hadoop, Hive, Airflow, Docker, Kubernetes, PostgreSQL, AWS, Git

Communication: English, Burmese, Chinese Mandarin, Japanese *(3 years of college-level coursework)*

Patents & Awards

Method for Transmitting Information to Device on a Network Based upon Information Associated with a Different Device on the Network, *Sivaramakrishnan, K.; Guan, D.; Kyaw T.; et al. 2016.* *US 20160125461 A1*

System and Method for Determining Related Digital Identities, *Sivaramakrishnan, K.; Guan, D.; Kyaw T.; et al. 2014.* *US 20140095320 A1*

Fourth Place, *National Matriculation Examination, Ministry of Education, Republic of the Union of Myanmar* *1997*